

A Preliminary Performance Model for Optimizing Software Packet Processing Pipelines

Ankit Bhardwaj, Atul Shree, Bhargav Reddy V, and Sorav Bansal

Department of Computer Science and Engineering, Indian Institute of Technology Delhi

Motivation and Objectives

Motivation: With increasing popularity of SDN, newer protocols are being developed over time. Researcher are developing DSLs to make the task easier. However, ease of programming doesn't guarantee application performance.

Goal: Develop a compiler that can automatically map a high-level specification to an underlying machine architecture in an optimal way.

Compilation Phases

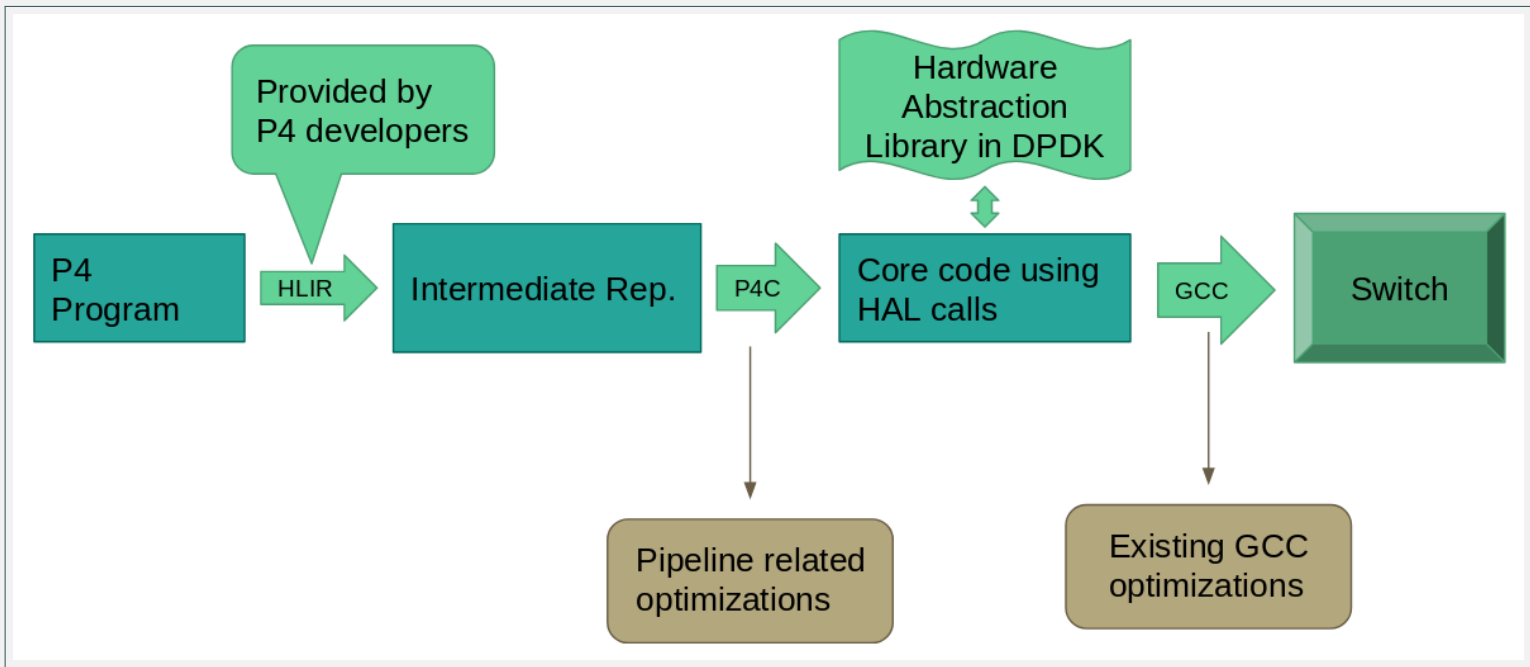


Figure: Compilation phases

Applications written in P4[1] are given as the input to P4C compiler and the compiler generated DPDK[2] based Applications. DPDK application is compiled using gcc to generate the target binary.

Packet Processing Pipeline

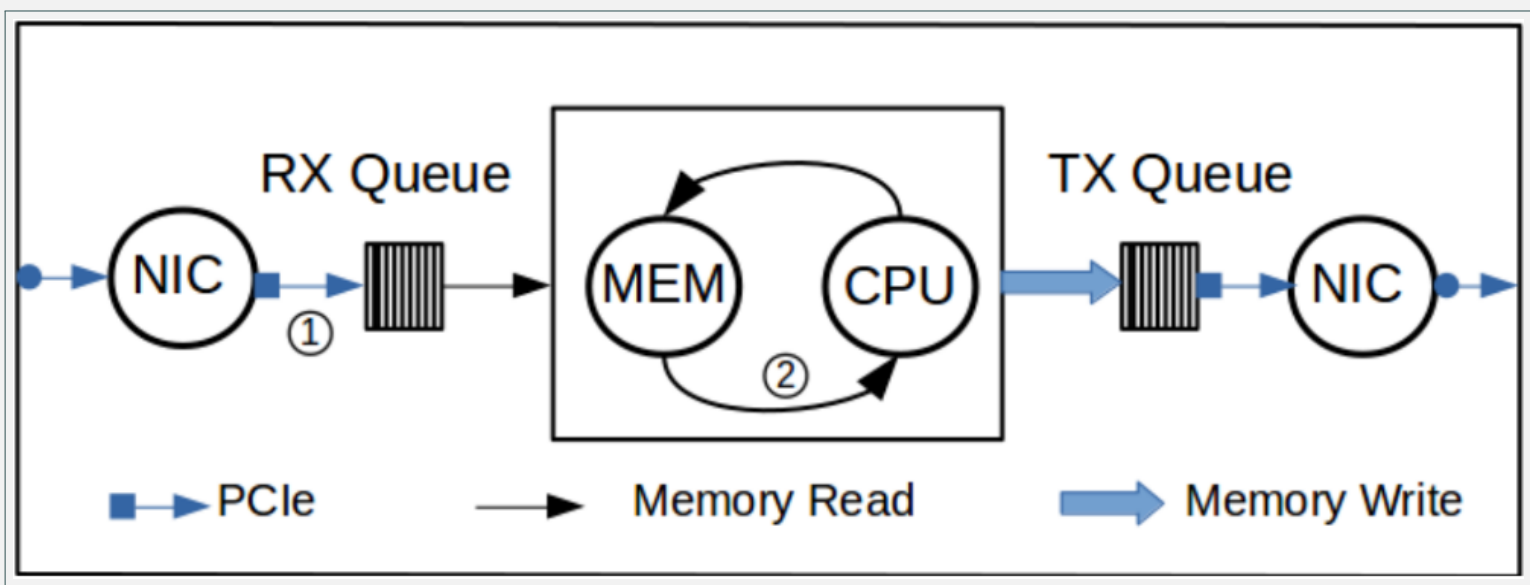


Figure: Packet Processing Pipeline

- * Exploit DMA bandwidth between NIC and Main Memory labeled ① in Figure
- * Exploit Memory Level Parallelism between CPU and Memory labeled ② in Figure

Loop-Fission for Batching, Sub-Batching, and Prefetching

```
sub app {
  for (i = 0; i < B; i++) {
    p = read_from_input_NIC();
    p = process_packet(p);
    write_to_output_NIC(p);
  }
}

sub app {
  for (i = 0; i < B; i++)
    p[i] = read_from_input_NIC();
  for (i = 0; i < B; i++)
    p[i] = process_packet(p[i]);
  for (i = 0; i < B; i++)
    write_to_output_NIC(p[i]);
}

sub process_packet(p) {
  for (i = 0; i < B; i++) {
    t1 = lookup_table1(p[i]);
    t2 = lookup_table2(p[i], t1);
    ...
  }
}

sub process_packet(p) {
  for (i = 0; i < B; i+=b) {
    for (j = i; j < i+b; j++)
      t1[j-i] = lookup_table1(p[j]);
    for (j = i; j < i+b; j++)
      t1[j-i] = lookup_table2(p[j], t1[j-i]);
    ...
  }
}

sub app {
  for (i = 0; i < B; i++)
    p[i] = read_from_input_NIC();
  for (i = 0; i < B; i++)
    p[i] = process_packet(p[i]);
  for (i = 0; i < B; i++)
    write_to_output_NIC(p[i]);
}
```

Performance Model

IO Subsystem, Memory Subsystem and CPU work in a pipeline fashion. Let the service rate of CPU, CPU-Memory, I/O-Memory DMA interface be c, m, d . *Throughput* of the application will be $\min(c, m, d)$. The value of b will vary with change in the parameters related to memory $f_{mem}(m, b)$. DMA interface is not linearly scalable and rate increases based on some function $f_{dma}(d, B)$.

$$Throughput = \min(c, f_{mem}(m, b), f_{dma}(d, B))$$

Experiments and Results




Figure: Effect of Batching and Prefetching Figure: Sensitivity of Throughput to B. Solid line $b = B$ and dotted lines $b = 1$. Figure: Sensitivity of Throughput to b. Batch Size $B = 128$ for these experiments.

Conclusion

We are using Batching(B) to exploit available parallelism between NIC and Memory interface, and prefetching to exploit the memory level parallelism between CPU and Memory interface. Sub-Batching(b) is being used to adjust the prefetch distance to take full advantage from software prefetching. The model to get the optimal value of b and B is preliminary and not fully automatic. We are running the applications to get the value of b and B and feeding the compiler to generate the optimal code. These optimizations can be used for wide variety of network and streaming applications.

Future Work

- Get more understanding of NIC-Memory interface and CPU-Memory interface.
- Make the model robust to be used for various kind of applications.
- Use the model for various architecture and do the appropriate changes.

References

[1] P4: Programming protocol-independent packet processors. *SIGCOMM Comput. Commun. Rev.*, 44(3):87–95, July 2014.

[2] Intel Data Plane Development Kit. <http://dpdk.org/>.

